

Tackling the Gradient Issues in Generative Adversarial Networks

Yanran Li

The Hong Kong Polytechnic University
yanranli.summer@gmail.com

Content

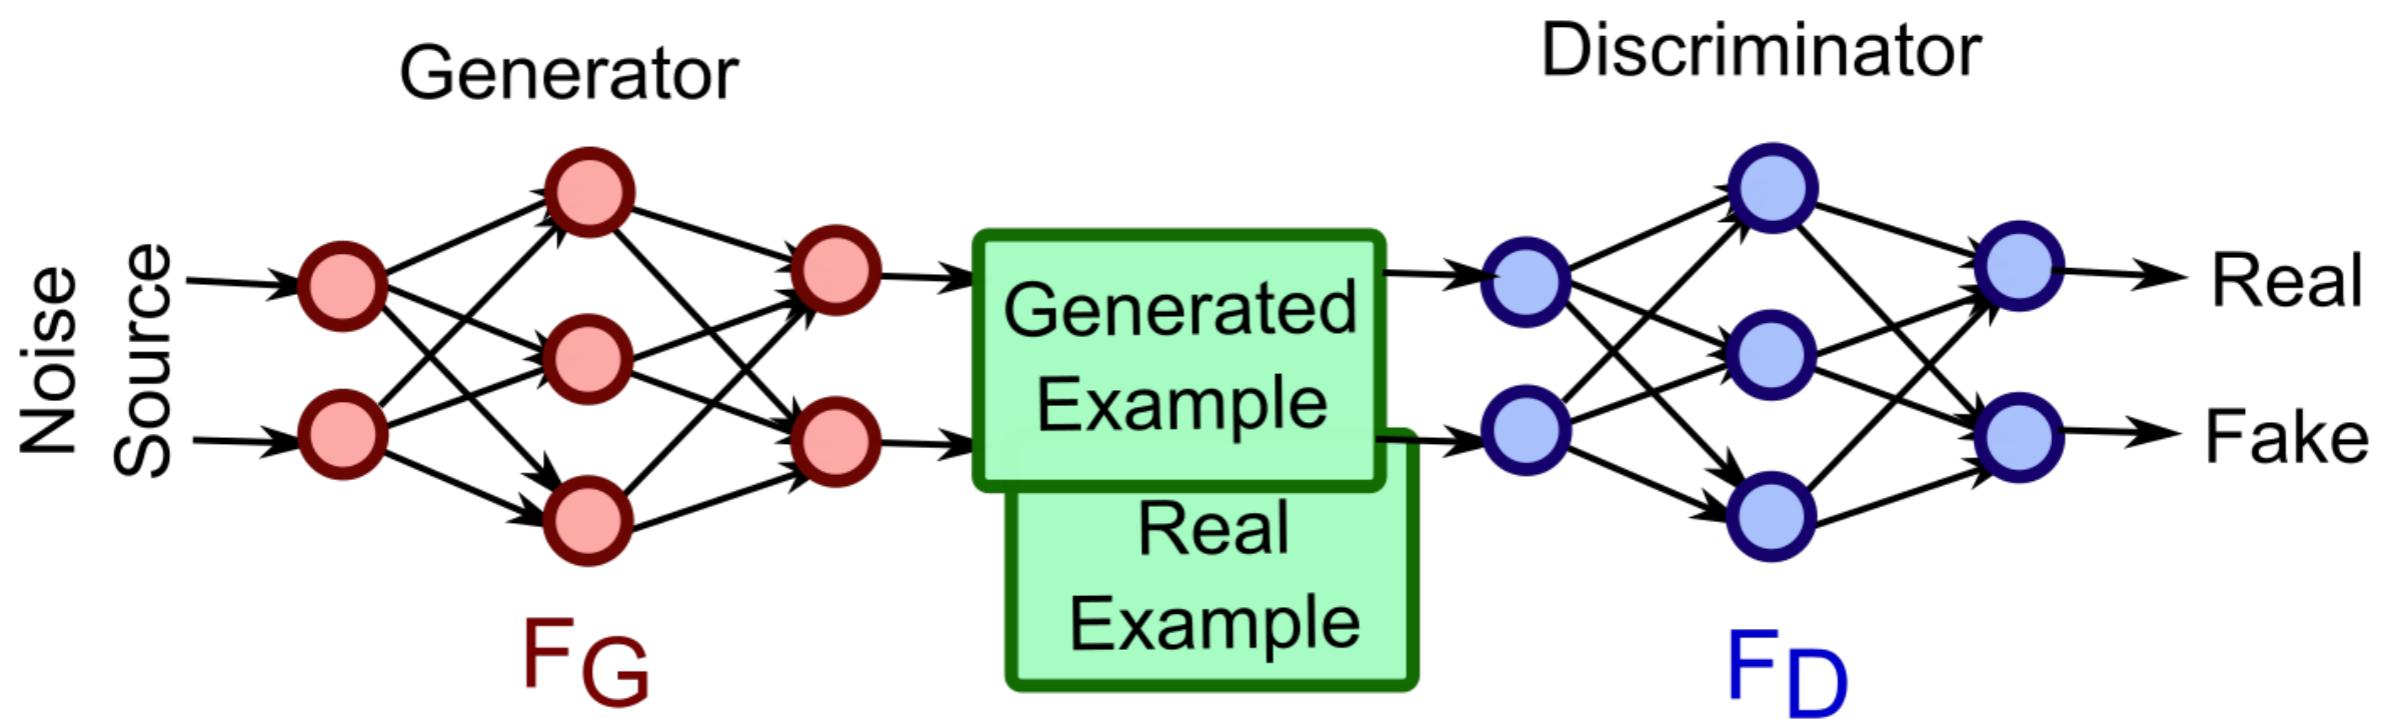
- Generative Adversarial Networks
 - Basics
 - Difficulties
- Solution 1: Encoder-incorporated
 - Mode Regularized GANs
 - Energy-based GANs, InfoGAN, etc.
 - *Noisy Input
- Solution 2: Wasserstein Distance
 - Wasserstein GANs and Improved Training of Wasserstein GANs

Content

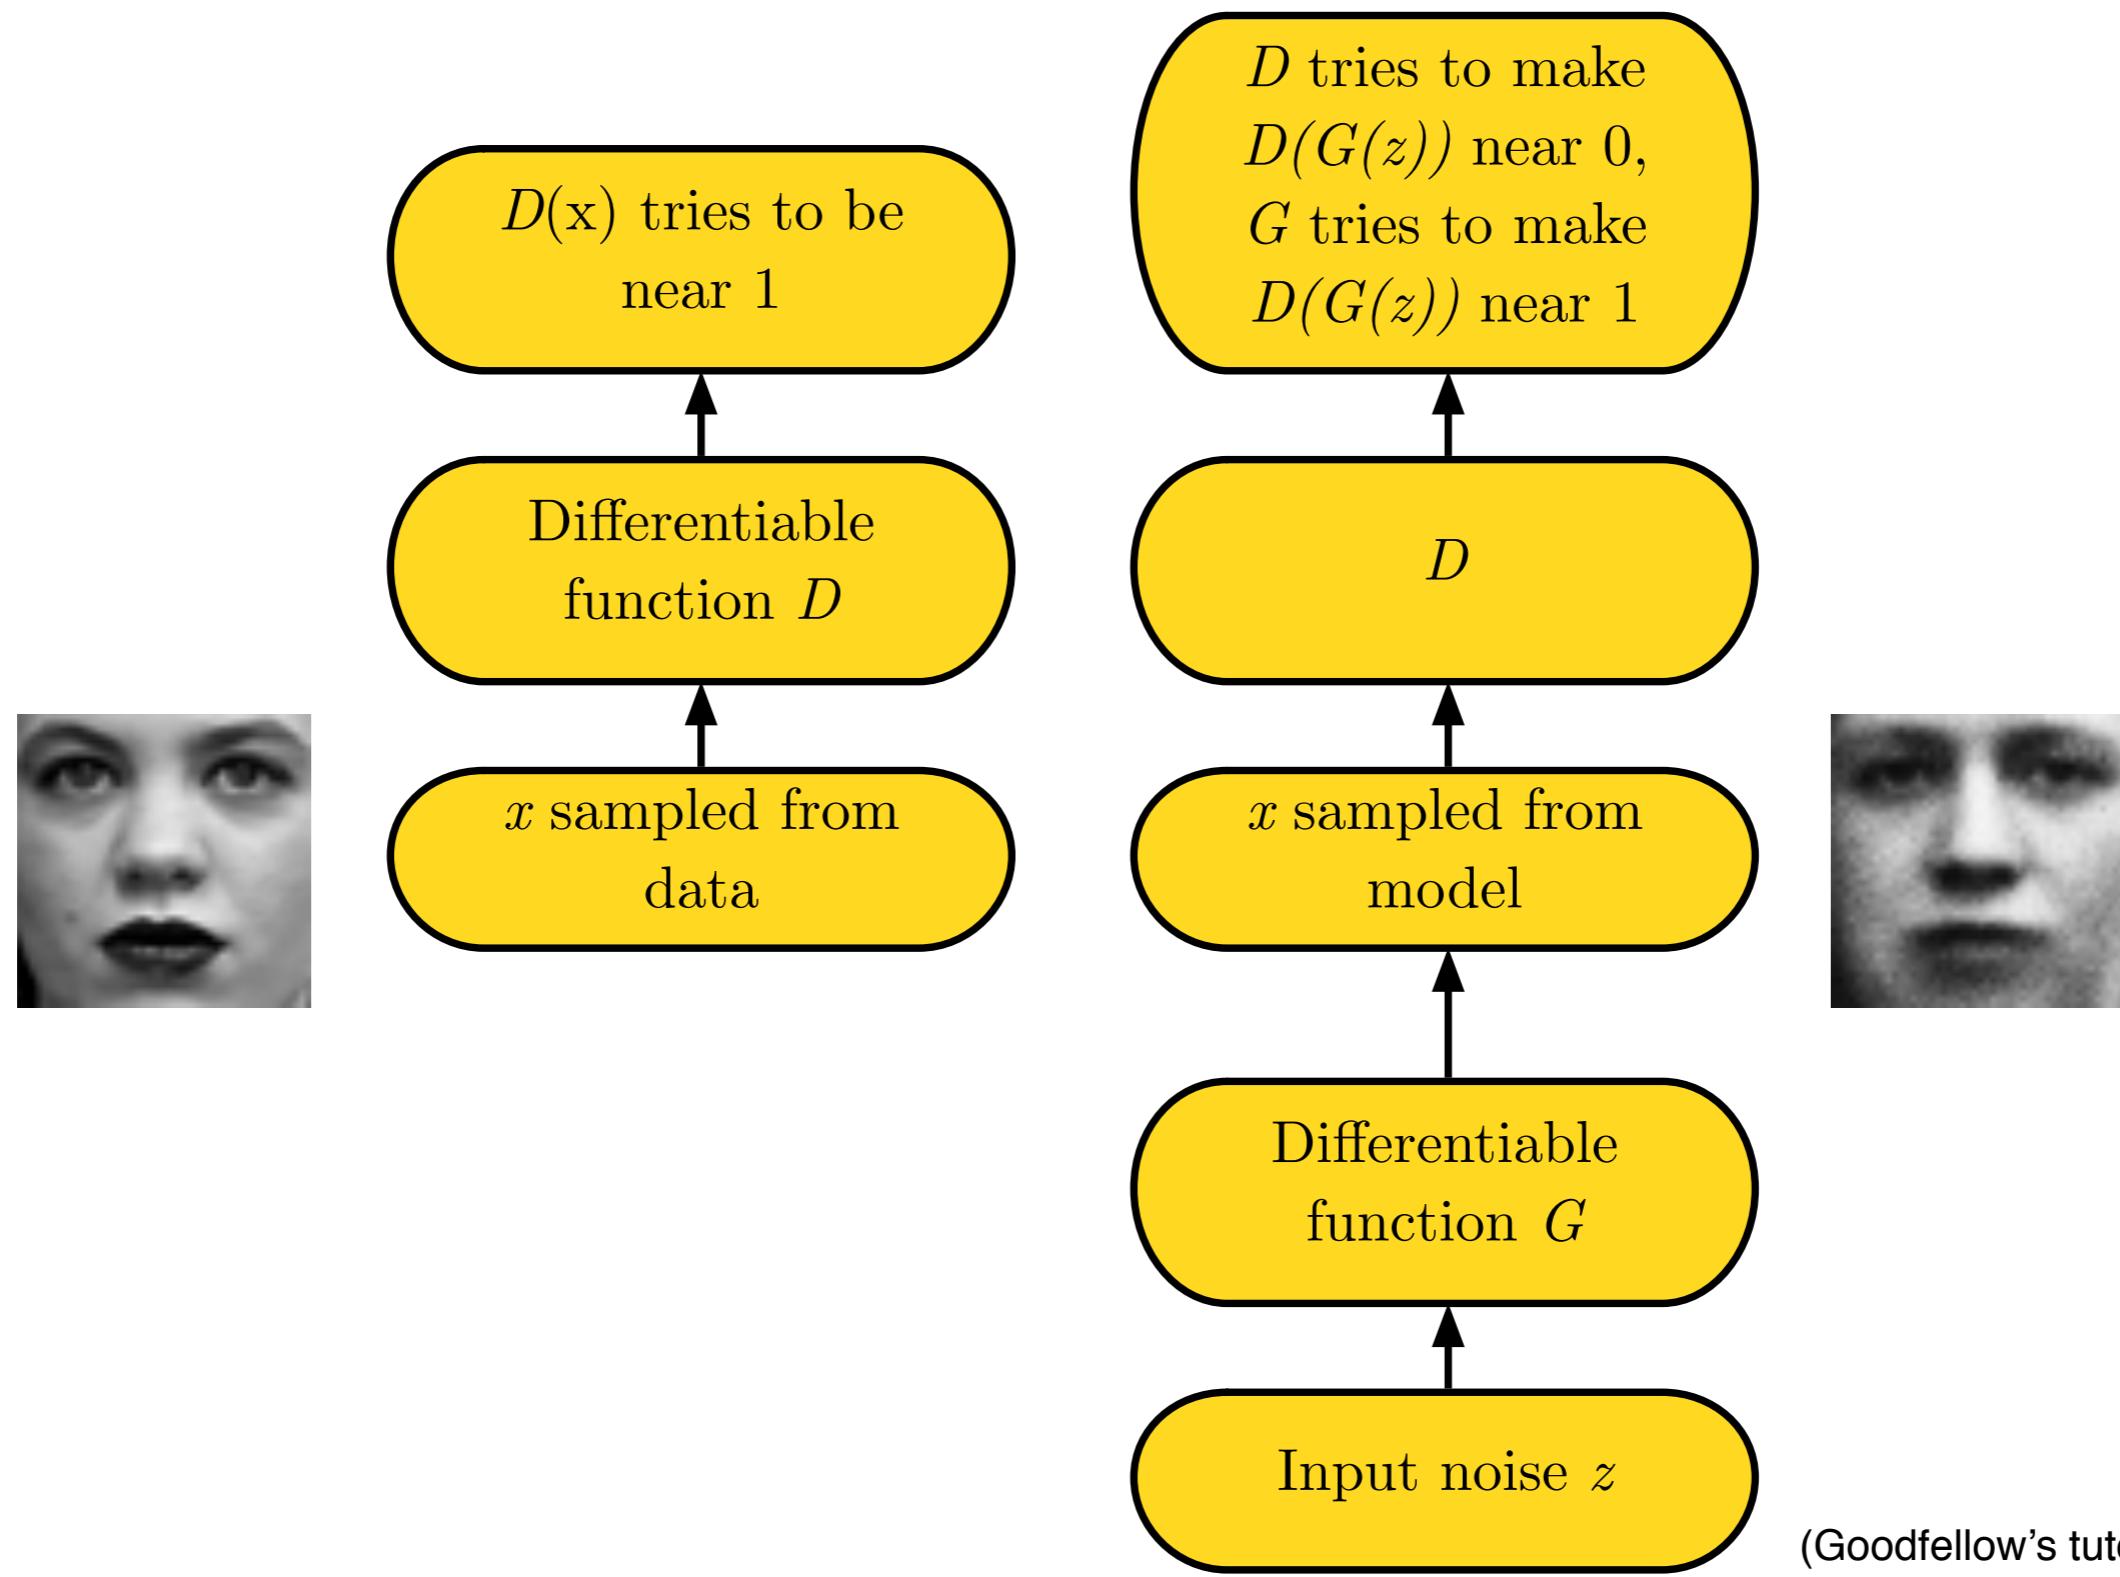
- **Generative Adversarial Networks**
 - Basics
 - Difficulties
- Solution 1: Encoder-incorporated
 - Mode Regularized GANs
 - Energy-based GANs, InfoGAN, etc.
 - *Noisy Input
- Solution 2: Wasserstein Distance
 - Wasserstein GANs and Improved Training of Wasserstein GANs

Generative Adversarial Networks

- A min-max game between two components: a generator \mathbf{G} and a discriminator \mathbf{D}



GANs Framework



(Goodfellow's tutorial)

Objectives for GAN

- The objective of D :

$$L(D, g_\theta) = \mathbb{E}_{x \sim \mathbb{P}_r} [\log D(x)] + \mathbb{E}_{x \sim \mathbb{P}_g} [\log(1 - D(x))]$$

- The objective of G :

- the original: $\mathbb{E}_{z \sim p(z)} [\log(1 - D(g_\theta(z)))]$

- the alternative: $\mathbb{E}_{z \sim p(z)} [-\log D(g_\theta(z))]$

- *Why alternative?*

Difficulty 1

- using the original form of the objective of \mathbf{G}

$$\mathbb{E}_{z \sim p(z)} [\log(1 - D(g_\theta(z)))]$$

will result in gradient vanishing issue of \mathbf{D} for \mathbf{G} because *intuitively*, at the very early phase of training, \mathbf{D} is very easy to be confident in detecting \mathbf{G} , so \mathbf{D} will output almost always 0

Difficulty 1

- using the original form of the objective of \mathbf{G}

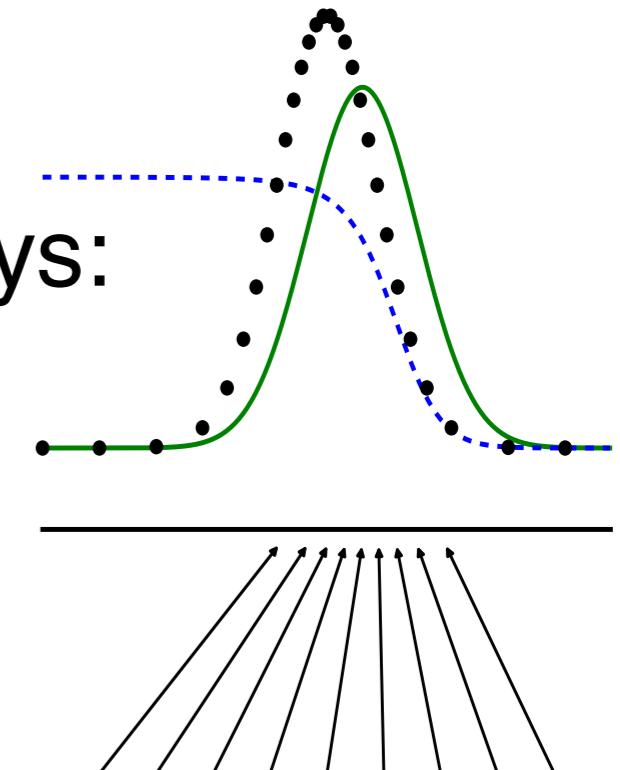
$$\mathbb{E}_{z \sim p(z)} [\log(1 - D(g_\theta(z)))]$$

will result in gradient vanishing issue of \mathbf{D} for \mathbf{G} because *theoretically*, when \mathbf{D} is *optimal*, minimizing the loss is equal to minimizing the *JS divergence* (Arjovsky & Bottou, 2017)

Difficulty 1

- The optimal D for any P_r and P_g is always:

$$D^*(x) = \frac{P_r(x)}{P_r(x) + P_g(x)}$$



and that

$$L(D^*, g_\theta) = 2JSD(\mathbb{P}_r \parallel \mathbb{P}_g) - 2 \log 2$$

so, when D is *optimal*, minimizing the loss is equal to minimizing the *JS divergence* (Arjovsky & Bottou, 2017)

Difficulty 1

- when:

$$L(D^*, g_\theta) = 2JSD(\mathbb{P}_r \parallel \mathbb{P}_g) - 2 \log 2$$

- The JS divergence for the two distributions P_r and P_g is (almost) always $\log 2$ because P_r and P_g hardly can overlap (Arjovsky & Bottou, 2017)
- This results in vanishing gradient in theory!

The alternative objective

- The alternative objective of \mathbf{G} :

$$\mathbb{E}_{z \sim p(z)} [-\log D(g_\theta(z))]$$

- Instead of minimizing, let \mathbf{G} maximize the log-probability of the discriminator being mistaken
- It is heuristically motivated that generator can still learn even when discriminator successfully rejects all generator samples, but not theoretically guaranteed

Difficulty 2

- using the alternative form of the objective of \mathbf{G}

$$\mathbb{E}_{z \sim p(z)} [-\log D(g_\theta(z))]$$

will result in gradient unstable issue and mode missing problem because *theoretically*, when D is *optimal*, minimizing the loss is equal to **minimizing** the *KL divergence* meanwhile **maximizing** the *JS divergence* (Arjovsky & Bottou, 2017):

$$KL(\mathbb{P}_{g_\theta} \parallel \mathbb{P}_r) - 2JSD(\mathbb{P}_{g_\theta} \parallel \mathbb{P}_r)]$$

Difficulty 2

- minimizing the *KL divergence* meanwhile maximizing the *JS divergence* is crazy:

$$KL(\mathbb{P}_{g_\theta} \parallel \mathbb{P}_r) - 2JSD(\mathbb{P}_{g_\theta} \parallel \mathbb{P}_r)]$$

- which results in gradient unstable issue

Difficulty 2

- minimizing the *KL divergence* is biased:

$$KL(\mathbb{P}_{g_\theta} \parallel \mathbb{P}_r) - 2JSD(\mathbb{P}_{g_\theta} \parallel \mathbb{P}_r)$$

- because *KL divergence* is asymmetric, and thus it is not equally treated when **G** generates a unreal sample and when **G** fails to generate real sample
- Therefore, **G** will generate too many few-mode but real samples, a safer strategy

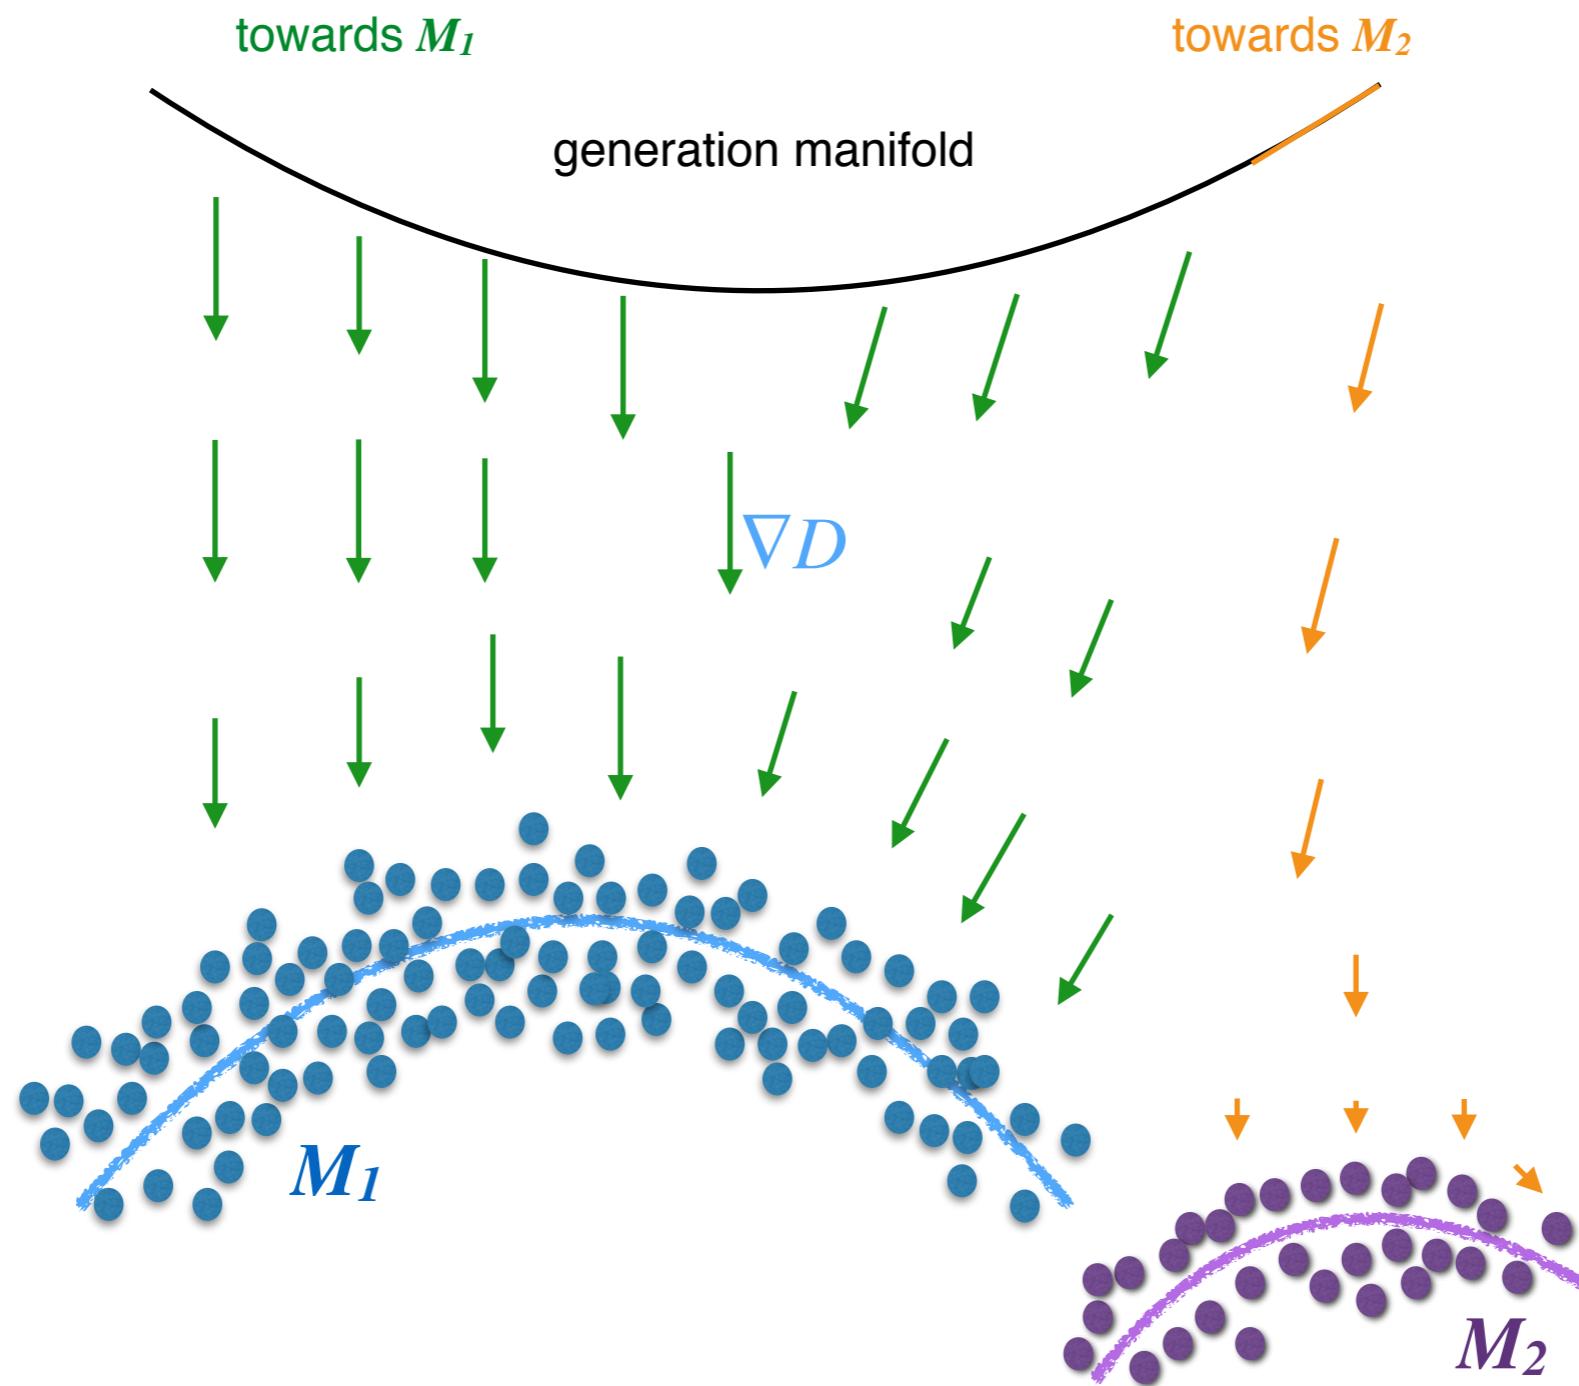
Content

- **Generative Adversarial Networks**
 - Basics
 - Difficulties
- **Solution 1: Encoder-incorporated**
 - Mode Regularized GANs
 - Energy-based GANs, Boundary Equilibrium GANs, etc.
 - *Noisy Input
- Solution 2: Wasserstein Distance
 - Wasserstein GANs and Improved Training of Wasserstein GANs

Solution 1: Encoder-incorporated

- Mode Regularized GANs (Che et al., 2017)
- Tackling the gradient vanishing issue and mode missing problem by incorporating an additional encoder E to:
 - (1) “enforce” P_r and P_g overlap
 - (2) “build a bridge” between *fake data* and *real data*

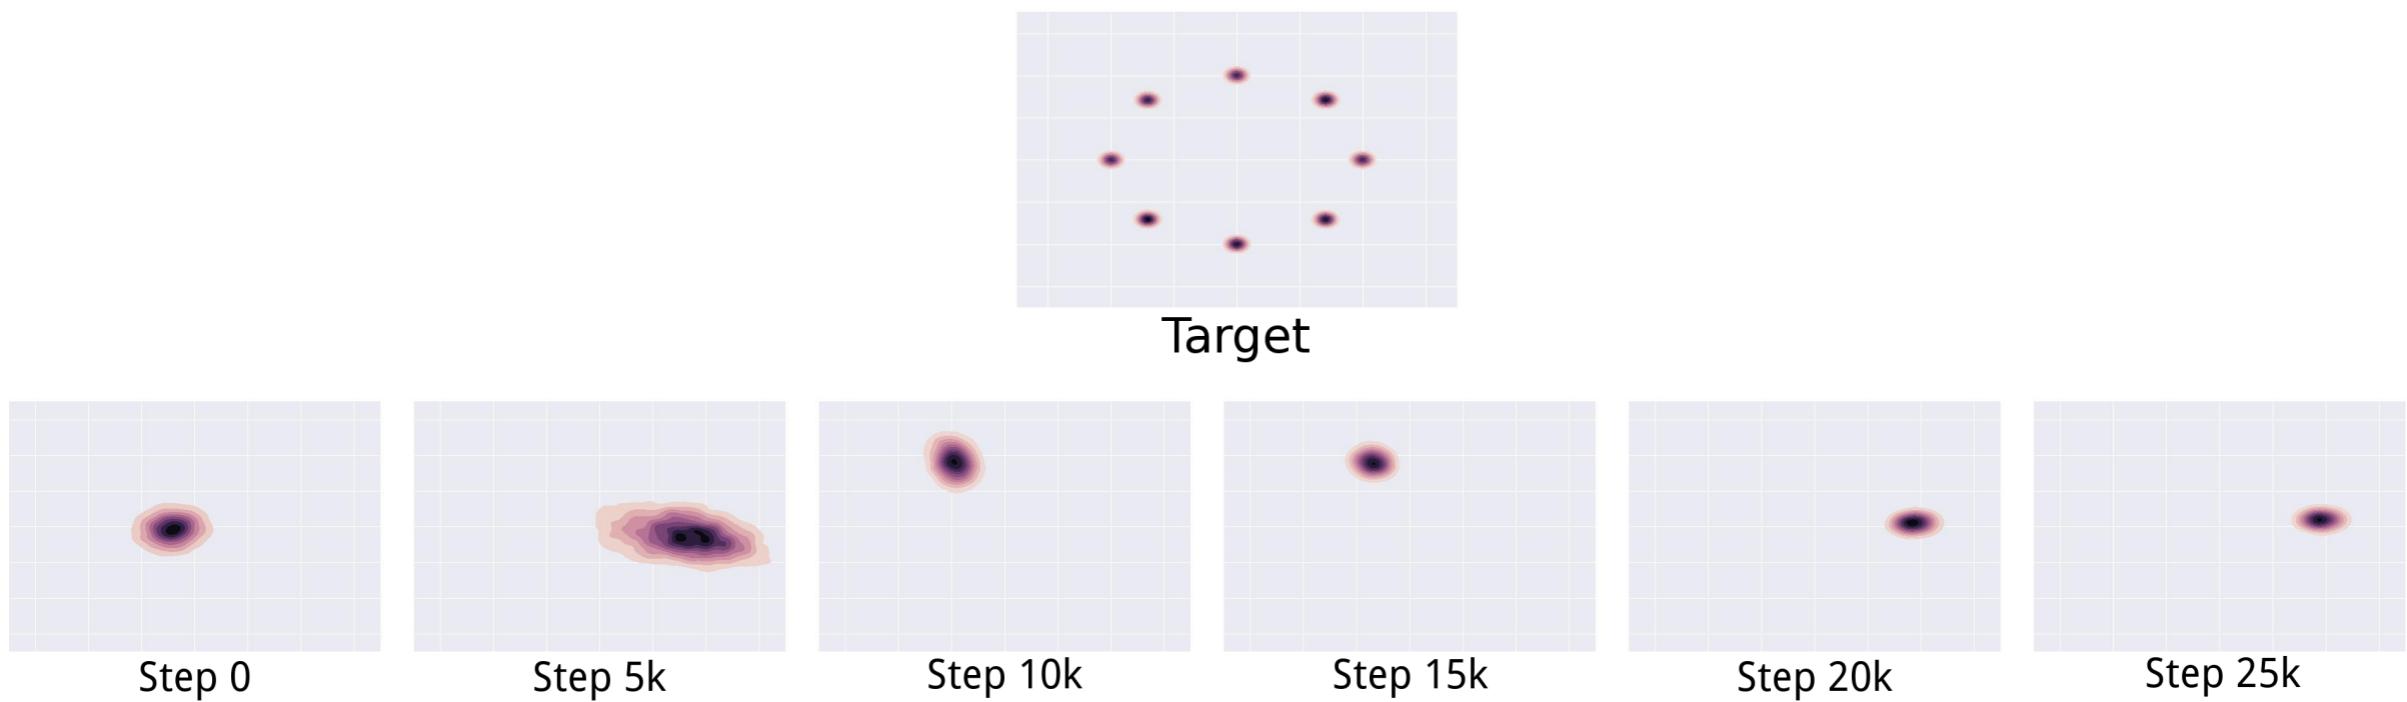
Mode Missing Problem



Mode Missing Problem

$$\min_G \max_D V(G, D) \neq \max_D \min_G V(G, D)$$

- **D** in inner loop: convergence to correct distribution
- **G** in inner loop: place all mass on most likely point



(Goodfellow's tutorial)
(Metz et al., 2016)

Mode Regularized GANs

- Regularized GANs
 - for encoder E : $\mathbb{E}_{x \sim p_d} [\lambda_1 d(x, G \circ E(x)) + \lambda_2 \log D(G \circ E(x))]$
 - for generator G :
- $-\mathbb{E}_z [\log D(G(z))] + \mathbb{E}_{x \sim p_d} [\lambda_1 d(x, G \circ E(x)) + \lambda_2 \log D(G \circ E(x))]$
- for discriminator D : same as vanilla GAN

Mode Regularized GANs

- Regularized GANs
 - for encoder E : $\mathbb{E}_{x \sim p_d} [\lambda_1 d(x, G \circ E(x)) + \lambda_2 \log D(G \circ E(x))]$
 - for generator G :
$$-\mathbb{E}_z [\log D(G(z))] + \mathbb{E}_{x \sim p_d} [\lambda_1 d(x, G \circ E(x)) + \lambda_2 \log D(G \circ E(x))]$$
 - for discriminator D : same as vanilla GAN
- But it still suffers from gradient vanishing!
- because D is still comparing between real data and fake data

Mode Regularized GANs

- Manifold-Diffusion GANs (MDGAN):

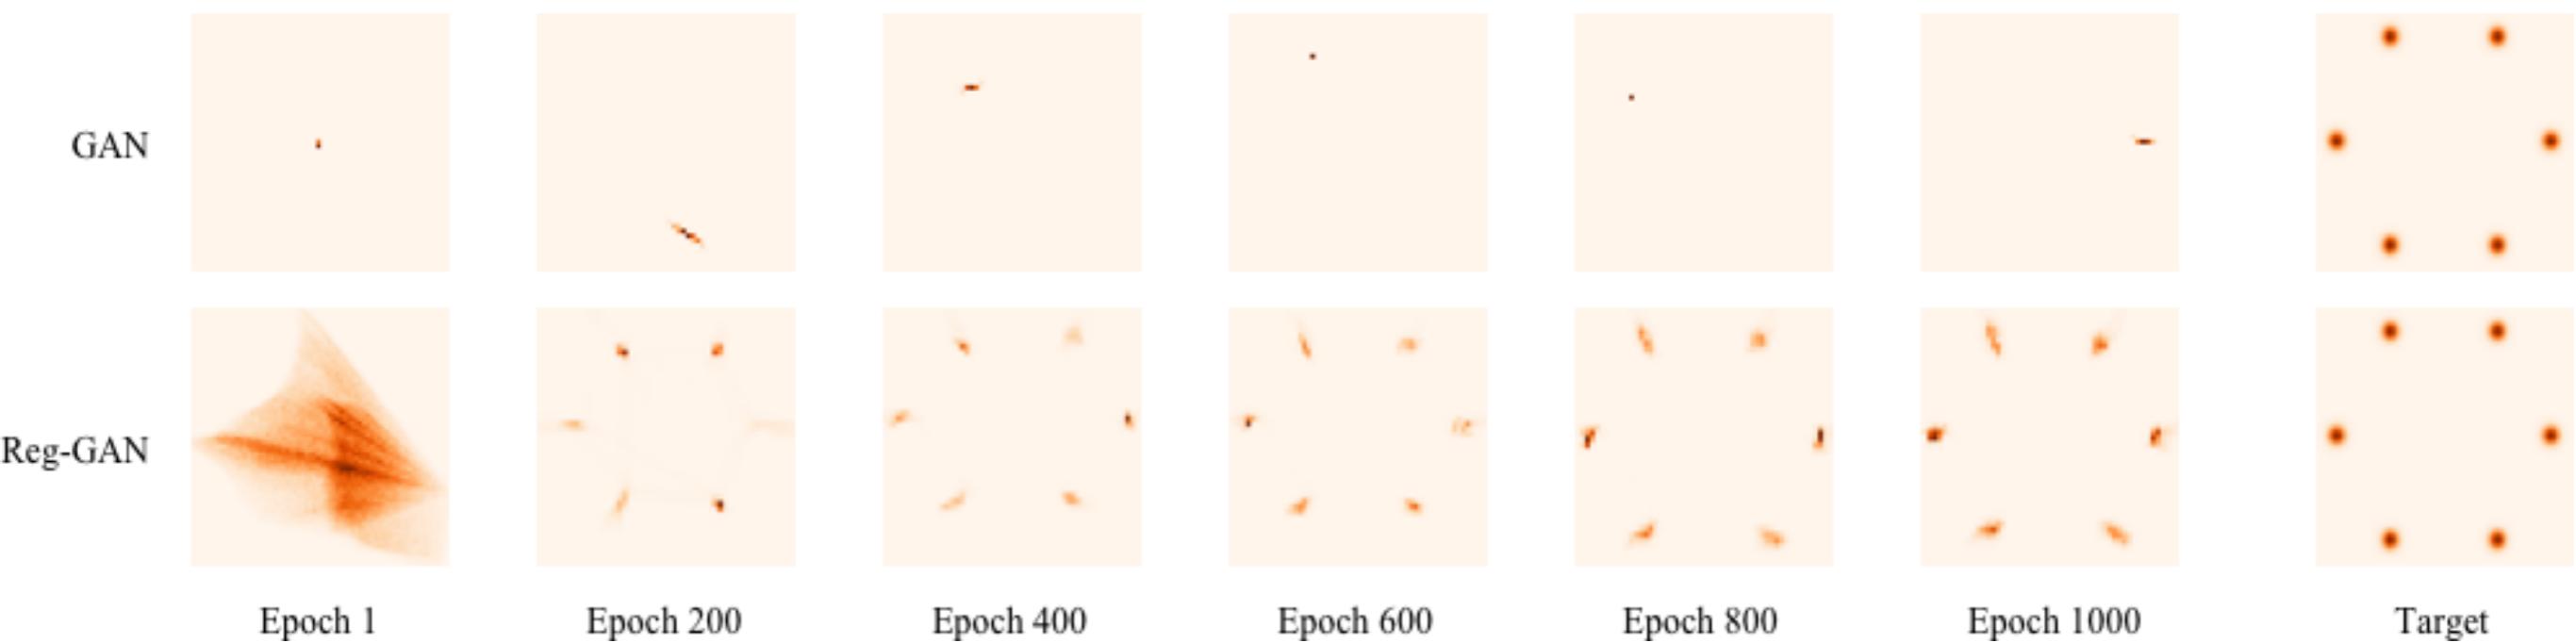
- for encoder E : $\mathbb{E}_{x \sim p_d} [\lambda_1 d(x, G \circ E(x)) + \lambda_2 \log D(G \circ E(x))]$
 - Manifold-step:
 - for generator G : $\lambda \log D_1(G(E(\mathbf{x}_i))) - \|\mathbf{x}_i - G(E(\mathbf{x}_i))\|^2$
 - for discriminator D : $\log D_1(\mathbf{x}_i) + \log(1 - D_1(G(E(\mathbf{x}_i))))$
 - Diffusion-step:
 - for generator G : $\log D_2(G(\mathbf{z}_i))$
 - for discriminator D : $\log D_2(G(E(\mathbf{x}_i))) + \log(1 - D_2(\mathbf{z}_i))$

Mode Regularized GANs

- Manifold-Diffusion GANs (MDGAN):

- for encoder E : $\mathbb{E}_{x \sim p_d} [\lambda_1 d(x, G \circ E(x)) + \lambda_2 \log D(G \circ E(x))]$
 - Manifold-step:
 - for generator G : $\lambda \log D_1(G(E(\mathbf{x}_i))) - \|\mathbf{x}_i - G(E(\mathbf{x}_i))\|^2$
 - for discriminator D : $\log D_1(\mathbf{x}_i) + \log(1 - D_1(G(E(\mathbf{x}_i))))$
 - Diffusion-step:
 - for generator G : $\log D_2(G(\mathbf{z}_i))$
 - for discriminator D : $\log D_2(G(E(\mathbf{x}_i))) + \log(1 - D_2(\mathbf{z}_i))$
- D is firstly comparing between real data and the encoded data — much harder!

Mode Regularized GANs



Mode Regularized GANs

MDGAN



Regularized
-GAN



ALI



VAEGAN
-trained



VAEGAN
-reported



DCGAN

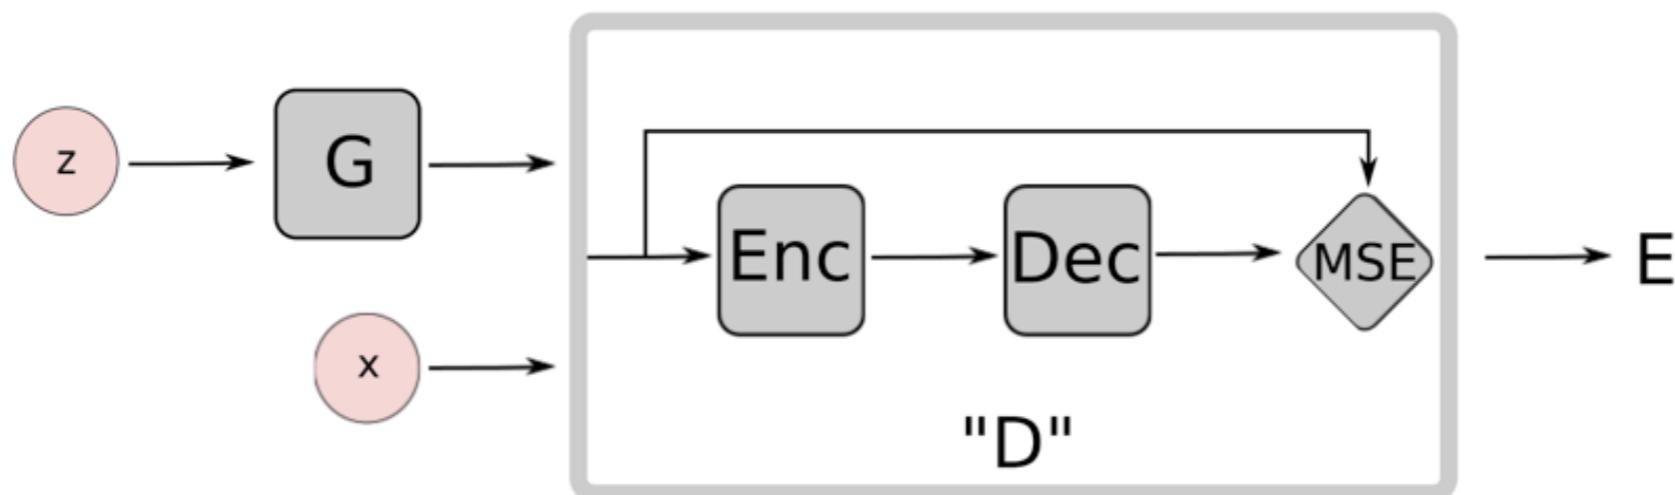


Solution 1: Encoder-incorporated

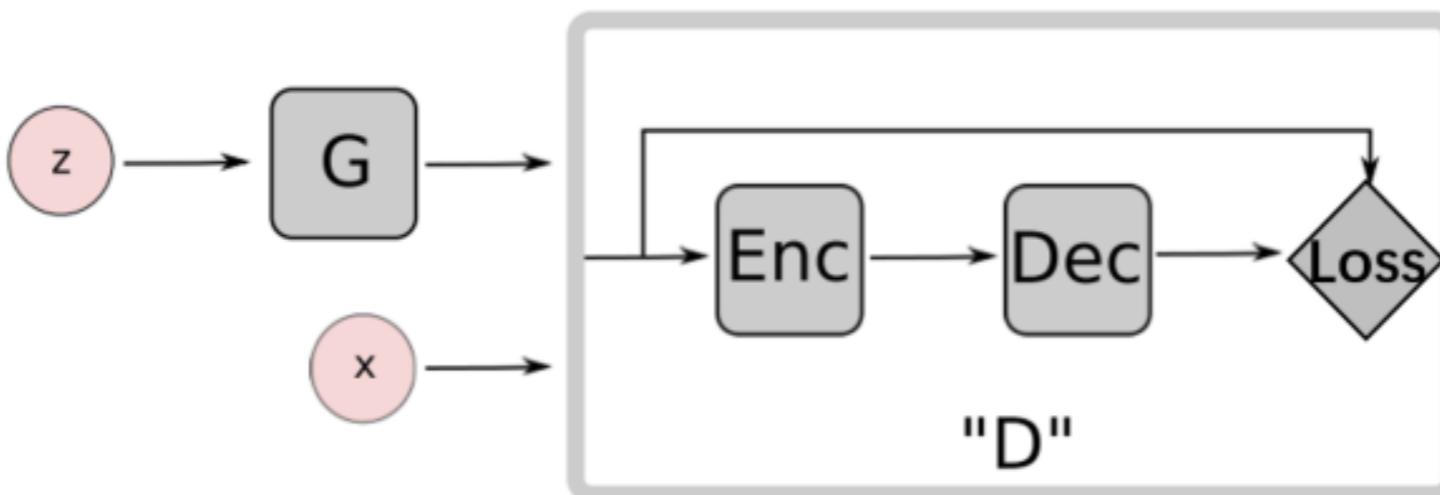
- Mode Regularized GANs (Che et al., 2017)
- Energy-based GANs (Zhao et al., 2017)
- Boundary Equilibrium GANs (Berthelot et al., 2017)
- etc.

Solution 1: Encoder-incorporated

- Energy-based GANs (Zhao et al., 2017)



- Boundary Equilibrium GANs (Berthelot et al., 2017)



Solution 1: **Noisy Input*

- Add noise to input (both real data and fake data) before passing into D (Arjovsky & Bottou, 2017)
- Add noise to layers in D and G (Zhao et al., 2017)
- Instance Noise (Sønderby et al., 2017)
- All these are indeed “enforcing” P_r and P_g to overlap

Content

- **Generative Adversarial Networks**
 - Basics
 - Difficulties
- Solution 1: Encoder-incorporated
 - Mode Regularized GANs
 - Energy-based GANs, Boundary Equilibrium GANs, etc.
 - *Noisy Input
- **Solution 2: Wasserstein Distance**
 - Wasserstein GANs

Solution 2: Wasserstein Distance

- Wasserstein GANs (Arjovsky et al., 2017)
- Wasserstein-1 Distance (Earth-Mover Distance):

$$W(\mathbb{P}_r, \mathbb{P}_g) = \inf_{\gamma \in \Pi(\mathbb{P}_r, \mathbb{P}_g)} \mathbb{E}_{(x,y) \sim \gamma} [\|x - y\|]$$

Solution 2: Wasserstein Distance

- Wasserstein GANs (Arjovsky et al., 2017)
- Wasserstein-1 Distance (Earth-Mover Distance):

$$W(\mathbb{P}_r, \mathbb{P}_g) = \inf_{\gamma \in \Pi(\mathbb{P}_r, \mathbb{P}_g)} \mathbb{E}_{(x,y) \sim \gamma} [\|x - y\|]$$

- *Why is it superior to KL and JS divergence?*

Solution 2: Wasserstein Distance

- Wasserstein-1 Distance (Earth-Mover Distance):

$$W(\mathbb{P}_r, \mathbb{P}_g) = \inf_{\gamma \in \Pi(\mathbb{P}_r, \mathbb{P}_g)} \mathbb{E}_{(x,y) \sim \gamma} [\|x - y\|]$$

where $\Pi(\mathbb{P}_r, \mathbb{P}_g)$ denotes the set of all joint distributions $\gamma(x, y)$ whose marginals are respectively \mathbb{P}_r and \mathbb{P}_g . Intuitively, $\gamma(x, y)$ indicates how much “mass” must be transported from x to y in order to transform the distributions \mathbb{P}_r into the distribution \mathbb{P}_g . The EM distance then is the “cost” of the optimal transport plan.

Solution 2: Wasserstein Distance

- Wasserstein-1 Distance (Earth-Mover Distance):

$$W(\mathbb{P}_r, \mathbb{P}_g) = \inf_{\gamma \in \Pi(\mathbb{P}_r, \mathbb{P}_g)} \mathbb{E}_{(x,y) \sim \gamma} [\|x - y\|]$$

- The distance is shown to have the desirable property that under mild assumptions
 - it is continuous everywhere and
 - differentiable almost everywhere.

Solution 2: Wasserstein Distance

- Wasserstein-1 Distance (Earth-Mover Distance):

$$W(\mathbb{P}_r, \mathbb{P}_g) = \inf_{\gamma \in \Pi(\mathbb{P}_r, \mathbb{P}_g)} \mathbb{E}_{(x,y) \sim \gamma} [\|x - y\|]$$

- The distance is shown to have the desirable property that under mild assumptions
 - And most importantly, it can reflect the distance of two distributions even if they do not overlap, and *thus can provide meaningful gradients*

Solution 2: Wasserstein Distance

- Wasserstein-1 Distance (Earth-Mover Distance):

$$W(\mathbb{P}_r, \mathbb{P}_g) = \inf_{\gamma \in \Pi(\mathbb{P}_r, \mathbb{P}_g)} \mathbb{E}_{(x,y) \sim \gamma} [\|x - y\|]$$

- By applying the Kantorovich-Rubinstein duality (Villani, 2008), Wasserstein GANs becomes:

$$\min_G \max_{D \in \mathcal{D}} \mathbb{E}_{\mathbf{x} \sim \mathbb{P}_r} [D(\mathbf{x})] - \mathbb{E}_{\tilde{\mathbf{x}} \sim \mathbb{P}_g} [D(\tilde{\mathbf{x}})]$$

Wasserstein GANs

- This new value function of WGAN gives rise to the additional requirement that the discriminator must lie within in the space of 1-Lipschitz functions:

$$\min_G \max_{D \in \mathcal{D}} \mathbb{E}_{\mathbf{x} \sim \mathbb{P}_r} [D(\mathbf{x})] - \mathbb{E}_{\tilde{\mathbf{x}} \sim \mathbb{P}_g} [D(\tilde{\mathbf{x}})]$$

- in other words, \mathcal{D} is the set of 1-Lipschitz functions
- To explain Lipschitz continuous is beyond today's topic

Wasserstein GANs

- This new value function of WGAN gives rise to the additional requirement that the discriminator must lie within in the space of 1-Lipschitz functions:

$$\min_G \max_{D \in \mathcal{D}} \mathbb{E}_{\mathbf{x} \sim \mathbb{P}_r} [D(\mathbf{x})] - \mathbb{E}_{\tilde{\mathbf{x}} \sim \mathbb{P}_g} [D(\tilde{\mathbf{x}})]$$

- To satisfy this requirement, WGAN enforces the weights of D lie within a compact space $[-c, c]$ by applying weight clipping

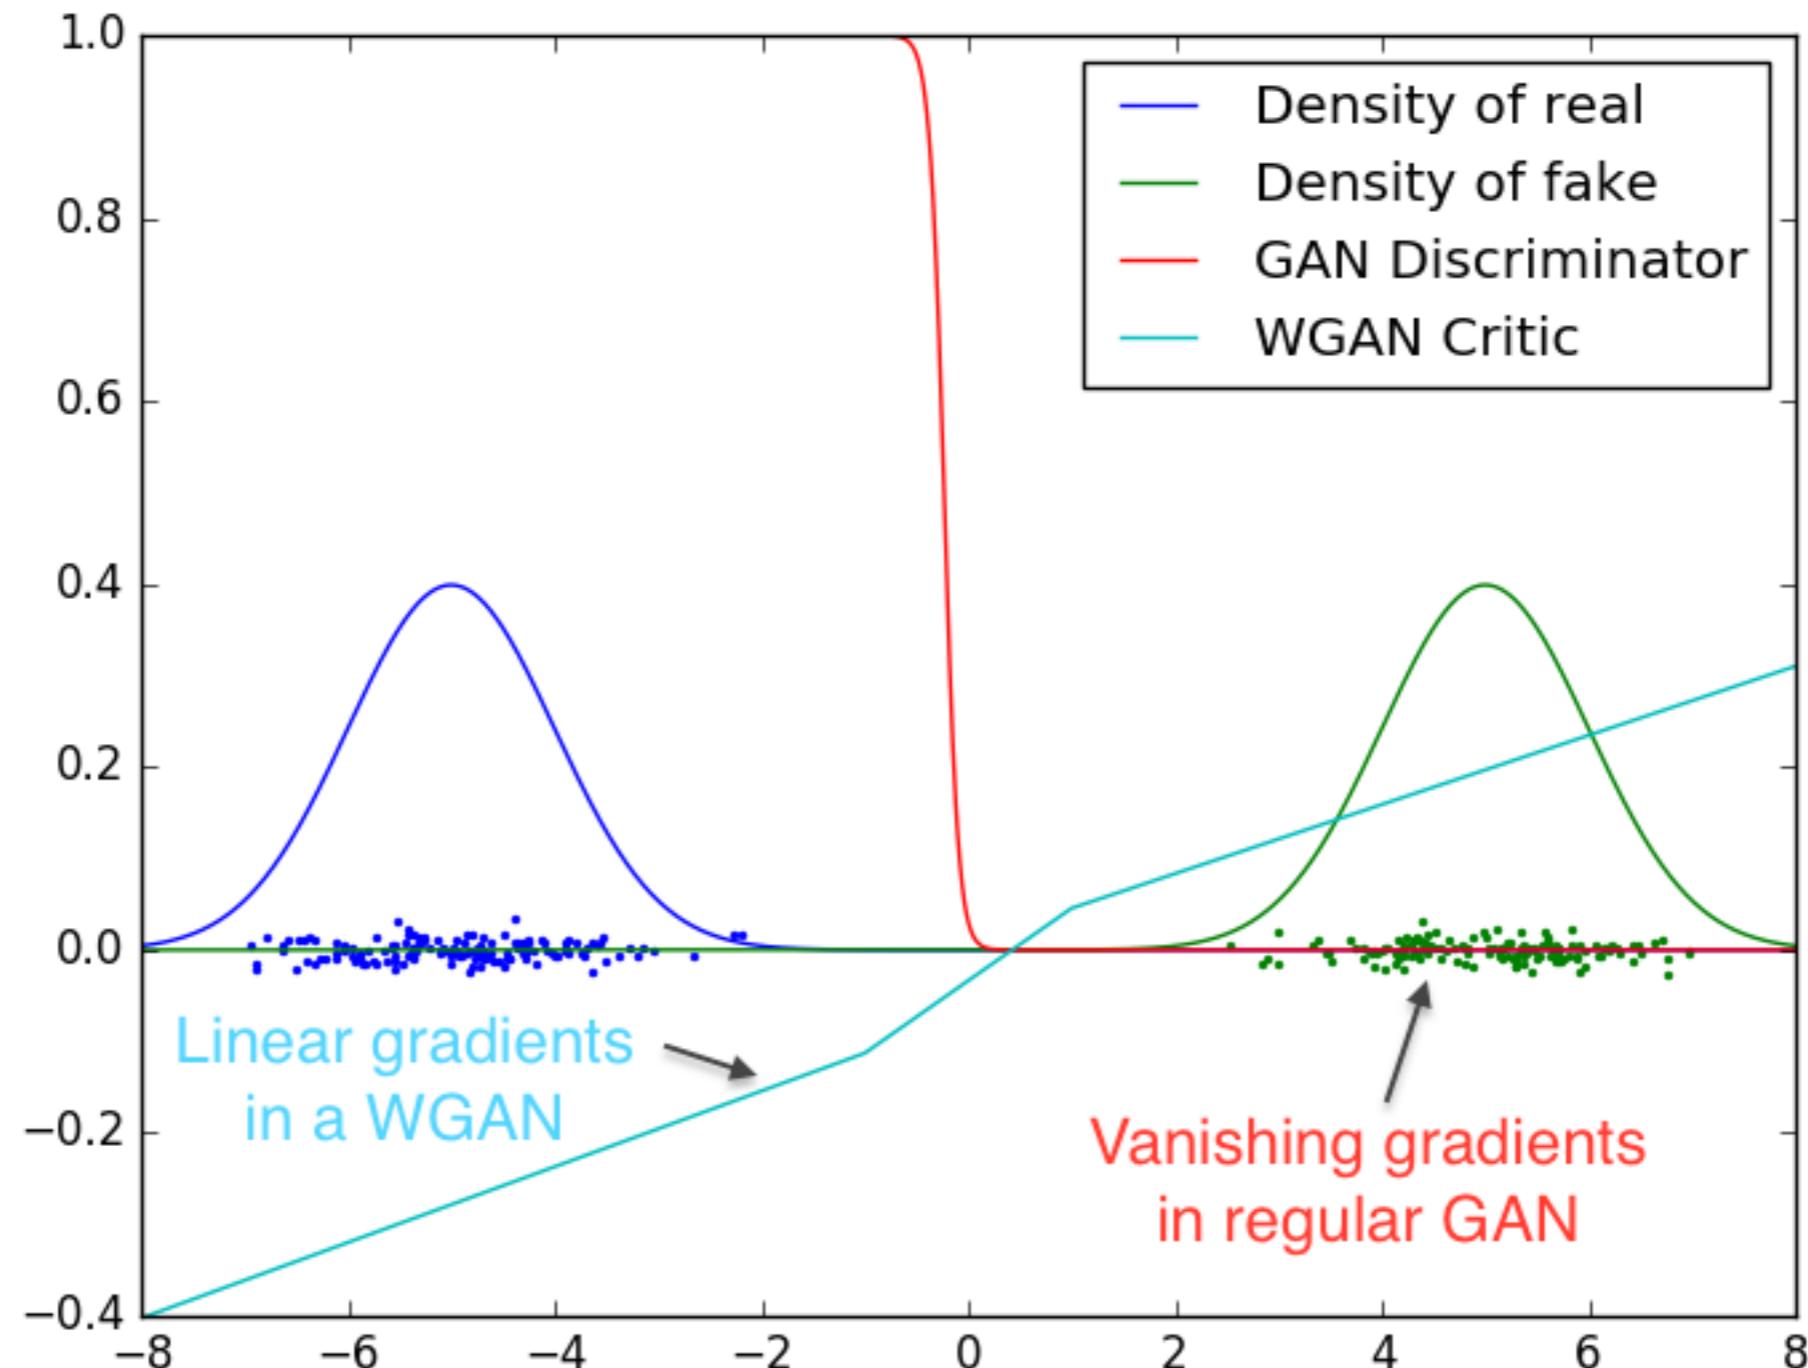
Wasserstein GANs

- This new value function of WGAN gives rise to the additional requirement that the discriminator must lie within in the space of 1-Lipschitz functions:

$$\min_G \max_{D \in \mathcal{D}} \mathbb{E}_{\mathbf{x} \sim \mathbb{P}_r} [D(\mathbf{x})] - \mathbb{E}_{\tilde{\mathbf{x}} \sim \mathbb{P}_g} [D(\tilde{\mathbf{x}})]$$

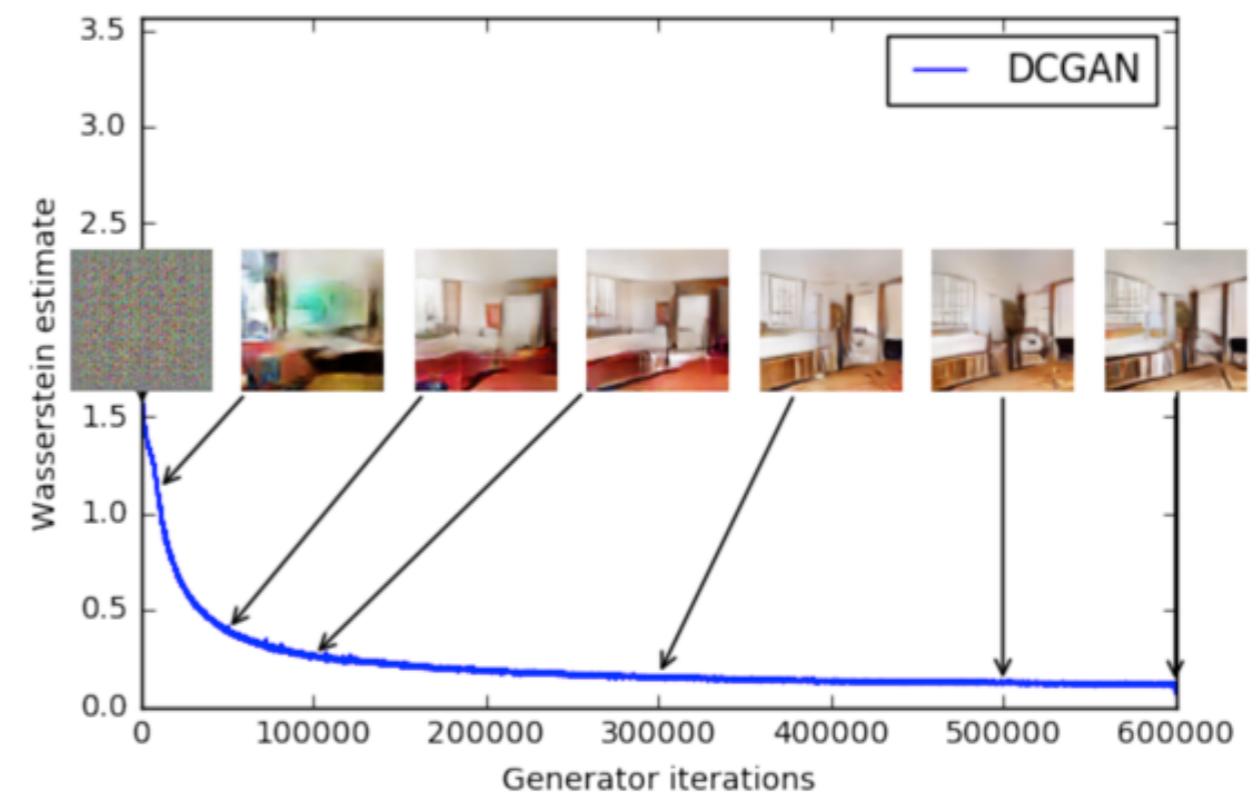
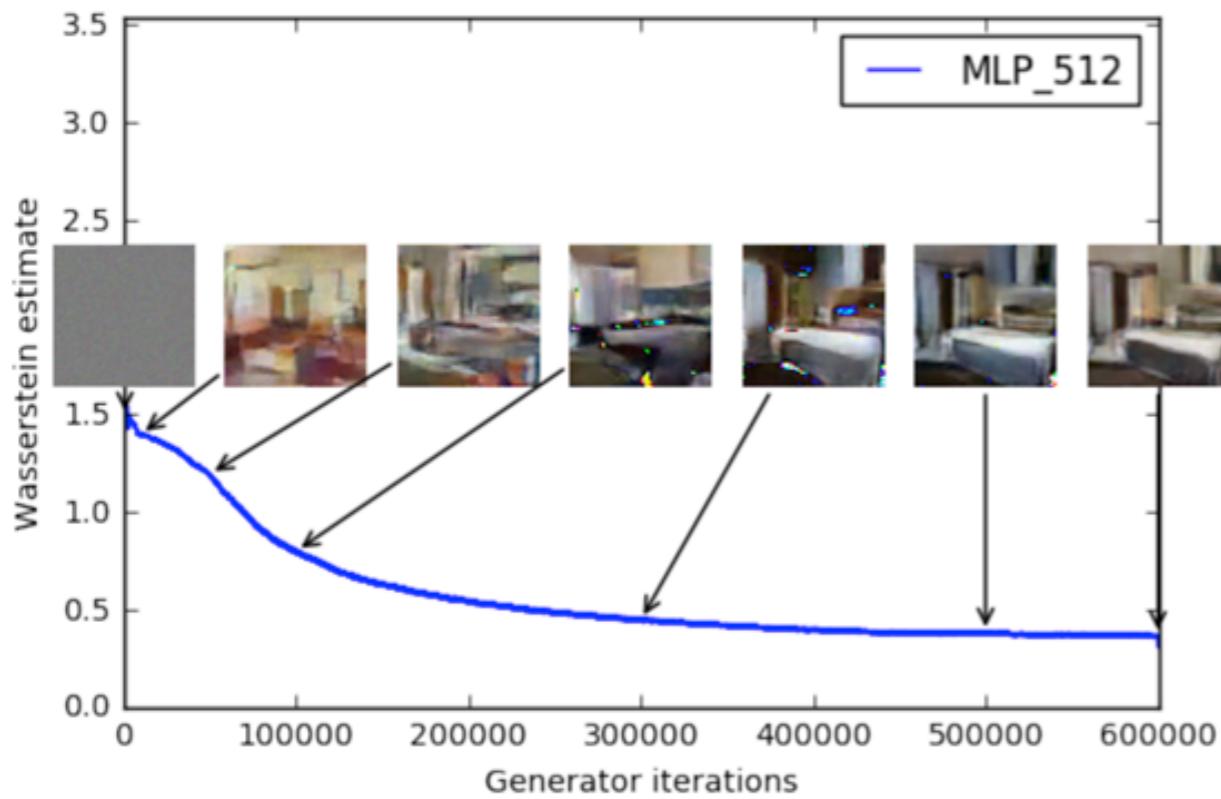
- Also, WGAN removes the sigmoid layer in \mathcal{D} because by using Wasserstein distance, \mathcal{D} in WGAN is doing regression rather than classification

Wasserstein GANs

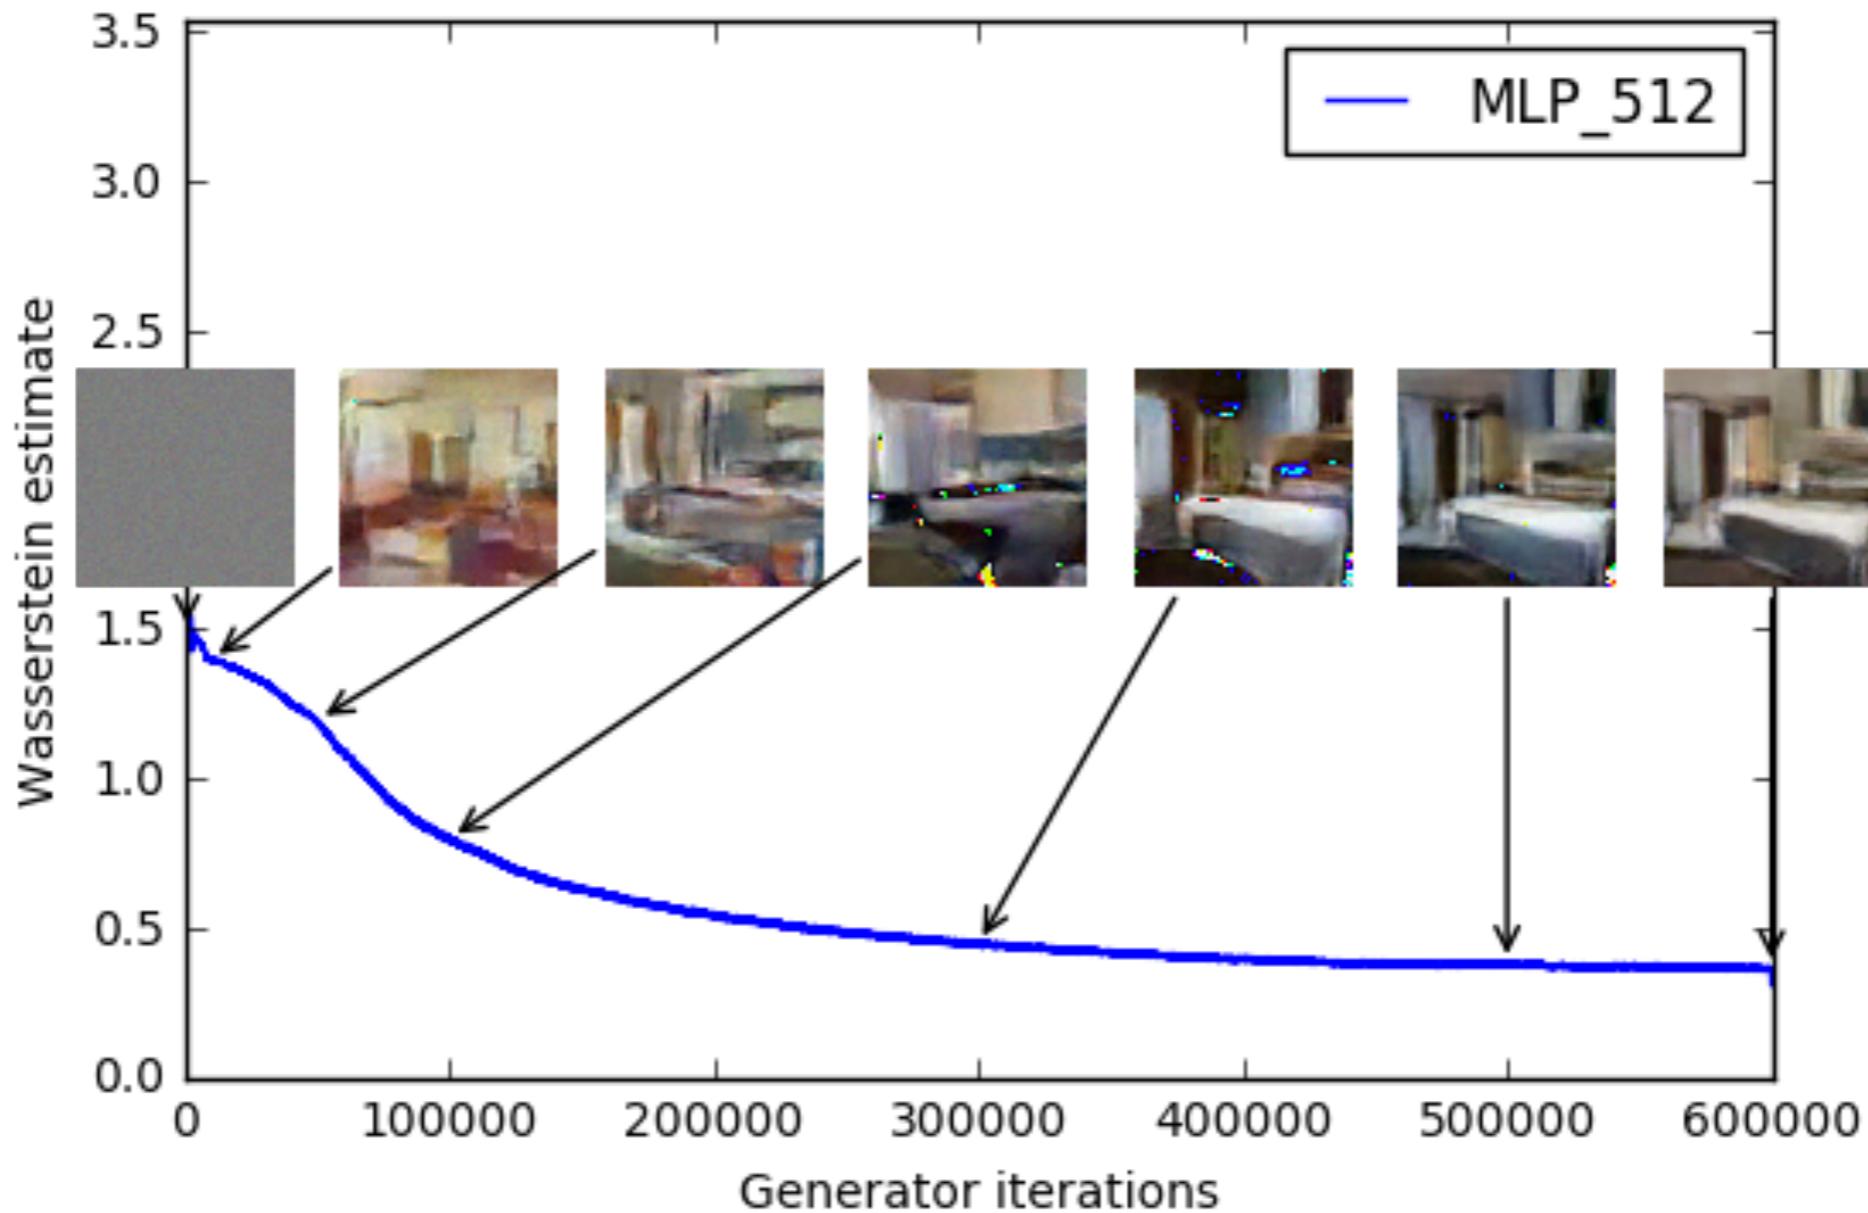


Wasserstein GANs

- This new value function of WGAN seems correlate with the quality of the generated samples:

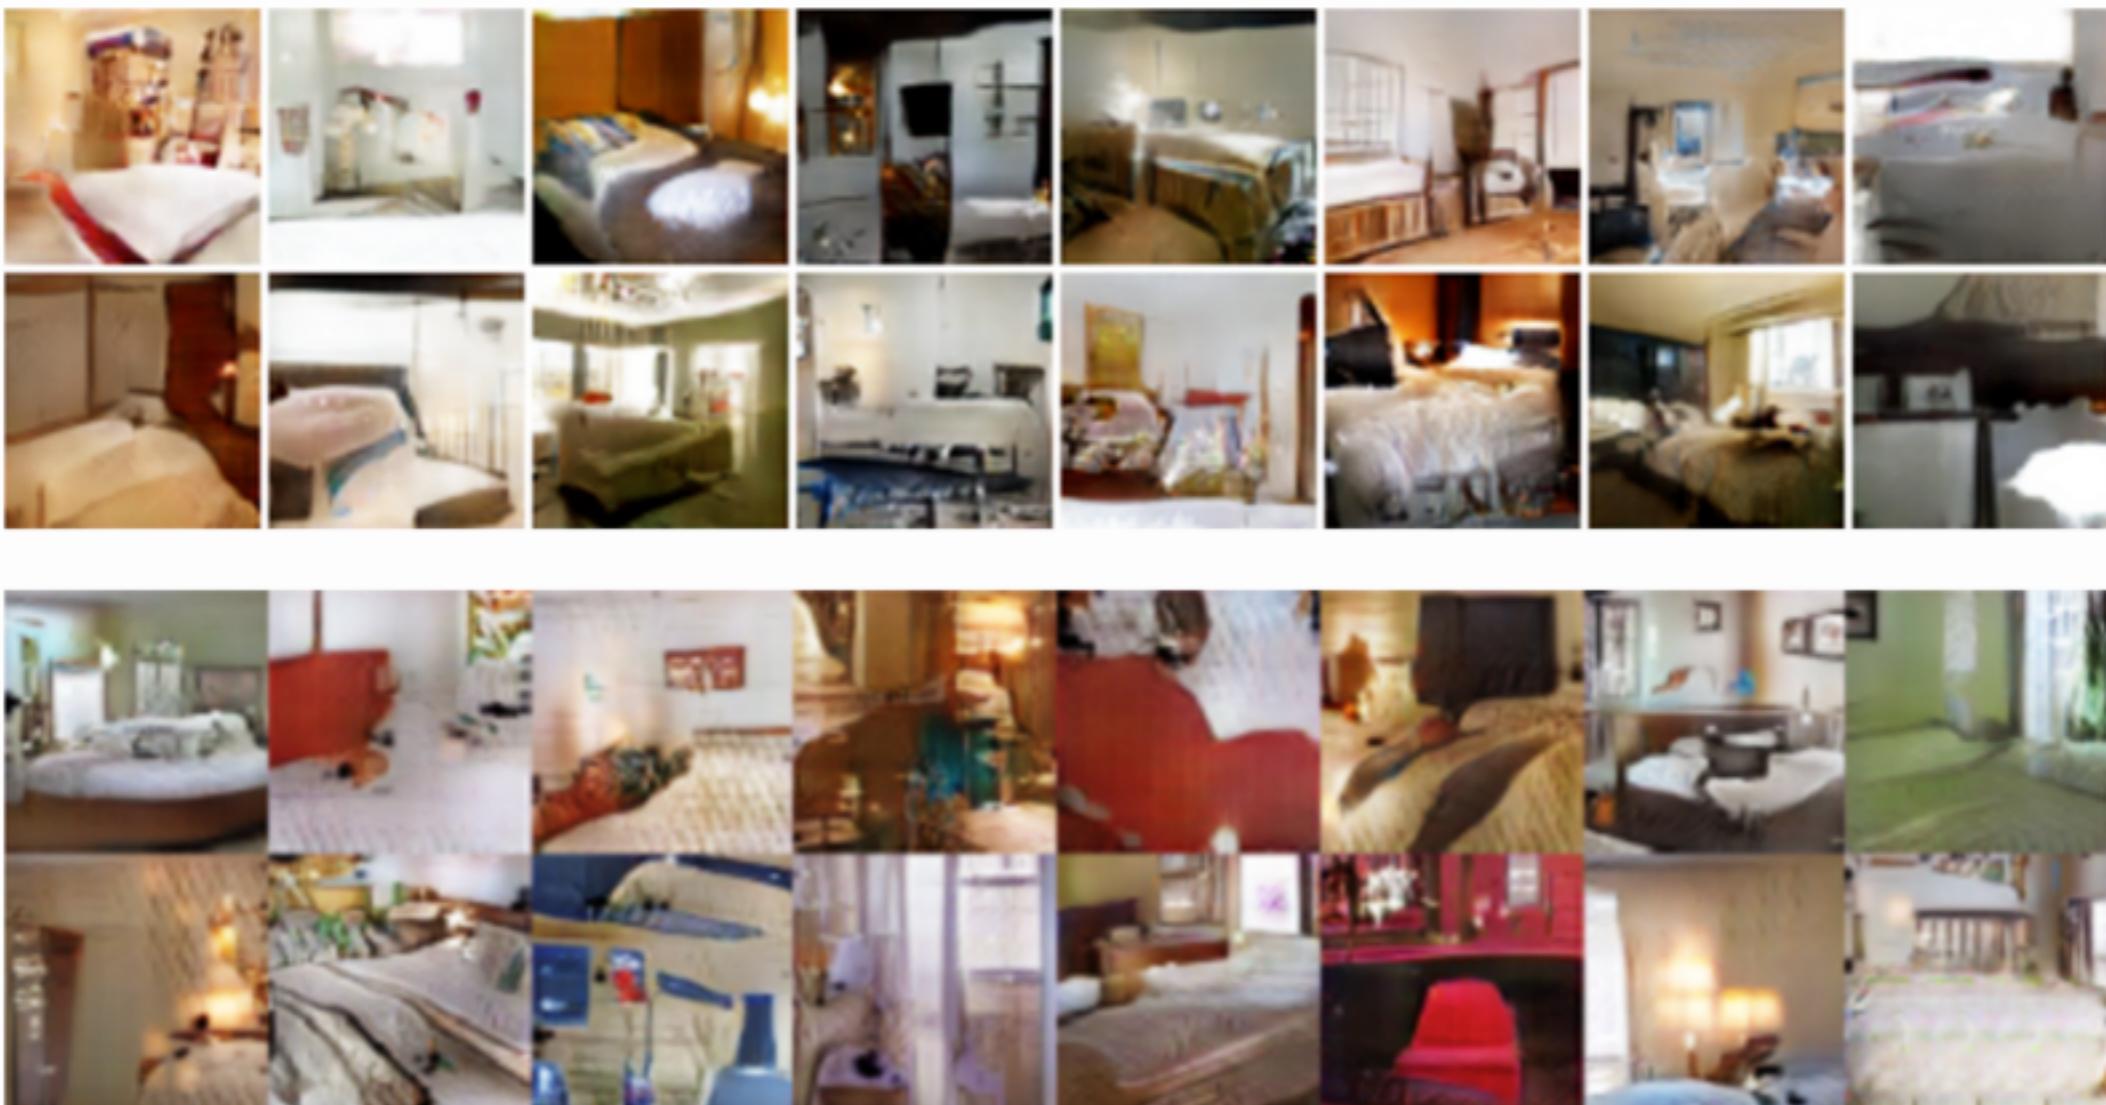


Wasserstein GANs



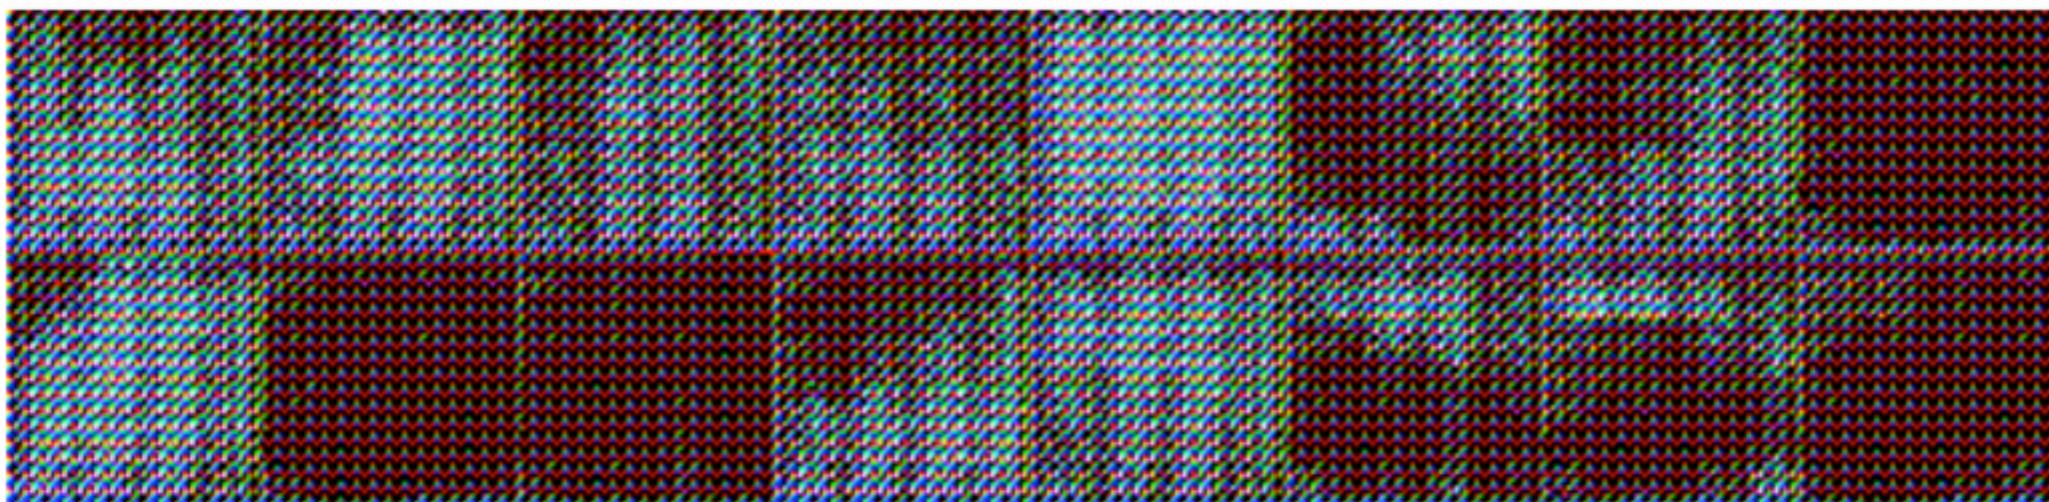
(Arjovsky et al., 2017)

Wasserstein GANs



Top: WGAN with the same DCGAN architecture. *Bottom:* DCGAN

Wasserstein GANs



Top: WGAN with DCGAN architecture, no batch norm. *Bottom:* DCGAN, no batch norm.

Wasserstein GANs



Top: WGAN with MLP architecture. *Bottom:* Standard GAN, same architecture.

Thanks for your attention!
Any questions?

References

- Arjovsky and Bottou, “Towards Principled Methods for Training Generative Adversarial Networks”. ICLR 2017.
- Goodfellow et al., “Generative Adversarial Networks”. ICLR 2014.
- Che et al., “Mode Regularized Generative Adversarial Networks”. ICLR 2017.
- Zhao et al., “Energy-based Generative Adversarial Networks”. ICLR 2017.
- Berthelot et al., “BEGAN: Boundary Equilibrium Generative Adversarial Networks”. arXiv preprint 2017.
- Sønderby, et al., “Amortised MAP Inference for Image Super-Resolution”. ICLR 2017.
- Arjovsky et al., “Wasserstein GANs”. arXiv preprint 2017.
- Villani, Cedric. “Optimal transport: old and new”, volume 338. Springer Science & Business Media, 2008